

Chapter 3

TrueSkill™ Rating System

TrueSkill [3] is a Gaussian rating system developed by Microsoft Research in 2007 to rank players on Xbox Live for the purpose of creating competitive matches amongst players with similar skill levels. While the Elo rating system is applicable to only two-player matches, the TrueSkill ranking system extends the use cases to multi-player matches. Section 3.1 describes Approximate Bayesian Inference — the mathematical backbone of the TrueSkill rating system and Section 3.2 discusses the prediction performances of several models built upon an open-source implementation of the TrueSkill rating system.

3.1 Approximate Bayesian Inference

The approach described in Section 2.2 can be modelled as a Bayesian Inference process. Define a vector $\mathbf{s} = [s_1, s_2, \dots, s_n]^\top$ representing the individual *skill* of n teams in a match (in the case of NCAA basketball tournament, $n = 2$). Further assume that each skill follows a Gaussian distribution $s_i \sim \mathcal{N}(\mu_i, \sigma_i^2)$ and collectively, the *prior* joint distribution of skills is a n-variate Gaussian distribution $p(\mathbf{s}) = \mathcal{N}_n(\mathbf{s}; \boldsymbol{\mu}, \boldsymbol{\Sigma})$.

In a match, each team is expected to exhibit a certain *performance* $p_i \sim \mathcal{N}(s_i, \beta_i^2)$ that varies around its skill. The match result can be represented as a set of rankings in ascending order for each team, namely, $\mathbf{r} = [r_1, r_2, \dots, r_n]^\top$ where $r_1 \leq r_2, \dots, r_{n-1} \leq r_n$. In reality, there are many factors that may affect the outcome of a match, but for the purpose of Bayesian Inference, the assumption is that it is the differences in team performances that cause the differences in team rankings. According to the Bayes' Theorem, the posterior distribution $p(\mathbf{s}|\mathbf{r})$ is given by

$$p(\mathbf{s}|\mathbf{r}) = \frac{P(\mathbf{r}|\mathbf{s})p(\mathbf{s})}{P(\mathbf{r})} = \frac{P(\mathbf{r}|\mathbf{s})p(\mathbf{s})}{\int P(\mathbf{r}|\mathbf{s})p(\mathbf{s})d\mathbf{s}} \quad (3.1)$$

The above formula demonstrates the use of Bayes' Theorem to calculate the updated skill distribution (posterior) given the match outcome (evidence) and the expected

tations (likelihood and prior) on the match outcome before the match. However, the posterior distribution is no longer a Gaussian distribution and oftentimes, the integral is intractable. A common solution is to calculate a Gaussian distribution $\mathcal{N}_n(\mathbf{s}; \hat{\boldsymbol{\mu}}, \hat{\boldsymbol{\Sigma}})$ as an approximation to the posterior distribution $p(\mathbf{s}|\mathbf{r})$ so that the Kullback-Leibler divergence between the two distributions is minimised.

A number of approximation techniques are available. For the TrueSkill rating system, it approximates the posterior distribution by setting up a *factor graph* which is a bipartite graph containing two kinds of vertices: variables and factors. A variable stores some value of interests and a factor represents some operations on one or more variables. The Bayesian Inference process is modelled as passing some *messages*, which are some real-valued functions of a variable or a factor, throughout the factor graph. Figure 3-1 shows an example factor graph of a simple two-player match. Based on skill s_i and performance p_i , it calculates the expected performance difference $p_1 - p_2$ and compares that with an externally supplied outcome d . The \mathbb{I} denotes an indicator function to assert whether d is consistent with the expectation, depending on which a different update is performed on s_1 and s_2 .




Figure 3-1: An example of factor graph

The ultimate goal of a factor graph is to obtain a set of update equations that indicate how each μ_i and σ_i should be updated. However, passing messages throughout

larger factor graphs can involve very intensive computations. As suggested by [1], there is an alternative but easier way of getting the same set of update equations based the following theorem.

Theorem 1. Let \mathbf{z} be a k-dimensional random vector $[z_1, z_2, \dots, z_k]^\top$ with a probability distribution function of the form

$$\frac{\phi_k(\mathbf{z})f(\mathbf{z})}{\int \phi_k(\mathbf{z})f(\mathbf{z})d\mathbf{z}} \quad (3.2)$$

where ϕ_k is a k-variate standard Gaussian distribution function. Then, the expectation of \mathbf{z} is given by

$$E[\mathbf{z}] = E\left[\frac{\nabla f(\mathbf{z})}{f(\mathbf{z})}\right] \quad (3.3)$$

and also

$$E[z_i z_j] = \delta_{ij} + E\left[\frac{\nabla^2 f(\mathbf{z})}{f(\mathbf{z})}\right] \quad (3.4)$$

where $\delta_{ij} = 1$ if $i = j$ and $\delta_{ij} = 0$ otherwise.

To link Equation 3.2 with the Bayes' Theorem in Equation 3.1, notice that the match outcome is dependent on performances and performances are related to skills; therefore, the likelihood probability $P(\mathbf{r}|\mathbf{s})$ can be expressed as some function of skill $f : \mathbf{s} \rightarrow [0, 1]$. Furthermore, since each $s_i \sim \mathcal{N}(\mu_i, \sigma_i^2)$, define the z-score vector $\mathbf{z} = [z_1, z_2, \dots, z_n]^\top$ where

$$z_i = \frac{s_i - \mu_i}{\sigma_i} \quad (3.5)$$

Therefore, the probability distribution $p(\mathbf{z}) = \phi_n(\mathbf{z})$ and the likelihood probability $P(\mathbf{r}|\mathbf{s})$ is also a function of \mathbf{z} . Then according to Equation 3.1,

$$p(\mathbf{z}|\mathbf{r}) = \frac{P(\mathbf{r}|\mathbf{z})p(\mathbf{z})}{P(\mathbf{r})} = \frac{\phi_n(\mathbf{z})f(\mathbf{z})}{\int \phi_n(\mathbf{z})f(\mathbf{z})d\mathbf{z}} \quad (3.6)$$

and therefore by Theorem 1, the values of interest μ_i^{new} and σ_i^{new} are related to $E[z_i]$ through

$$\begin{aligned} \mu_i^{new} &= E[s_i] = E[\mu_i + \sigma_i z_i] = \mu_i + \sigma_i E[z_i] \\ &= \mu_i + \sigma_i E\left[\frac{\partial f(\mathbf{z})/\partial z_i}{f(\mathbf{z})}\right] \end{aligned} \quad (3.7)$$

and

$$\begin{aligned}\sigma_i^{new} &= \text{Var}[s_i] = \sigma_i^2 \text{Var}[z_i] = \sigma_i^2 (E[z_i^2] - E[z_i]^2) \\ &= \sigma_i^2 \left(1 + E \left[\frac{\nabla^2 f(\mathbf{z})}{f(\mathbf{z})} \right]_{ii} - E \left[\frac{\partial f(\mathbf{z}) / \partial z_i}{f(\mathbf{z})} \right]^2 \right)\end{aligned}\quad (3.8)$$

The above update equations are applicable to the general cases where n teams are involved. However, for NCAA tournament matches that are between only two teams, a simpler version of the update equation can be derived. Now consider two teams with skill $\mathbf{s} = [s_i, s_j]^\top$. In general, there is no draw in basketball matches, but for completeness, assume the *draw margin* is ε such that team i beats team j when $s_i - s_j > \varepsilon$ and the outcome is considered as draw when $|s_i - s_j| \leq \varepsilon$. Furthermore, define $c^2 = 2\beta^2 + \sigma_i^2 + \sigma_j^2$. Then

$$f(\mathbf{z}) = \Phi\left(\frac{\mu_i - \mu_j - \varepsilon}{c}\right) \quad (3.9)$$

when team i beats team j . Alternatively,

$$f(\mathbf{z}) = \Phi\left(\frac{\varepsilon - (\mu_i - \mu_j)}{c}\right) - \Phi\left(\frac{-\varepsilon - (\mu_i - \mu_j)}{c}\right) \quad (3.10)$$

when team i draws with team j . In both cases, Φ represents the standard Gaussian cumulative distribution function.

Another approximation applied above is to assume $\mathbf{z} = \mathbf{0}$, namely, the two teams are assumed to have skills that are equal to their mean skills ($\mathbf{s} = \boldsymbol{\mu}$). The rationale for the above equations is to construct a Gaussian difference distribution [5] $\mathcal{N}_{i-j}(s_{i-j}; \mu_i - \mu_j, c^2)$ considering the performance variance β^2 . Then for the case that team i beats team j , it is to calculate the probability that $\mu_i - \mu_j - \varepsilon > 0$ using a standard Gaussian Φ function. Similar arguments apply to the draw case.

Now that $f(\mathbf{z})$ is known, for the case team i beats team j

$$\begin{aligned}\frac{\partial f(\mathbf{z})}{\partial z_i} &= \frac{\partial f(\mathbf{z})}{\partial \mu_i} \cdot \frac{\partial \mu_i}{\partial z_i} = \sigma_i \frac{\partial}{\partial \mu_i} \Phi\left(\frac{\mu_i - \mu_j - \varepsilon}{c}\right) \\ &= \frac{\sigma_i}{c} \phi\left(\frac{\mu_i - \mu_j - \varepsilon}{c}\right)\end{aligned}\quad (3.11)$$

where ϕ is the standard Gaussian probability density function. Following Equation 3.7

and 3.8 , the skill of team i can be updated accordingly:

$$\mu_i^{new} = \mu_i + \frac{\sigma_i^2}{c} \cdot \frac{\phi\left(\frac{\mu_i - \mu_j - \varepsilon}{c}\right)}{\Phi\left(\frac{\mu_i - \mu_j - \varepsilon}{c}\right)} \quad (3.12)$$

$$\sigma_i^{new} = \sigma_i^2 \left\{ 1 + \frac{\sigma_i^2}{c} \left[-\frac{\mu_i - \mu_j - \varepsilon}{c^2} \cdot \frac{\phi\left(\frac{\mu_i - \mu_j - \varepsilon}{c}\right)}{\Phi\left(\frac{\mu_i - \mu_j - \varepsilon}{c}\right)} - \frac{1}{c} \cdot \frac{\phi\left(\frac{\mu_i - \mu_j - \varepsilon}{c}\right)^2}{\Phi\left(\frac{\mu_i - \mu_j - \varepsilon}{c}\right)^2} \right] \right\} \quad (3.13)$$

Detailed derivations are omitted and can be found in [1]. Update equations for team j and for the case of draw can be derived in a similar way. By means of Approximate Bayesian Inference, the TrueSkill rating system is able to update the skill of each team upon the end of a match and to estimate a team's probability of winning an upcoming match.

3.2 Implementation